

Concurrency

COMP 2911, 18s2

Thread Basics

- A **thread** is a single sequence of execution within a process
- **Multi-threading** refers to multiple threads of control within a single process
- Process versus a thread
 - A process has a self-contained execution environment with its own dedicated memory space
 - Threads are light-weight compared to processes
 - Threads run within a single process, share the same address space and files of the process.
 - All Java programs have at least one thread, known as the main thread, which is created by the JVM at the program's start, when the `main()` method is invoked with the main thread.

Creating and Starting Threads

- Create a thread in Java as:

```
Thread thread = new Thread();
```

- To start the Java thread you will call its start() method, like this:

```
thread.start();
```
- The example above, does not specify any code for the thread to execute
- The thread will stop again right away after it is started
- There are two ways to specify what code a thread should execute
 - Extend the Thread Class
 - Implement the interface Runnable

Extending the Thread class

- Create a subclass of Thread and override the run() method
- The run() method is executed, when you invoke start() on the thread

```
public class MyThread extends Thread {  
    public void run() {  
        System.out.println("MyThread running");  
    }  
}
```

- To create and start the above thread

```
MyThread aThread = new MyThread();  
aThread.start();
```

Implementing the Runnable interface

- Create a class that implements `Runnable` and the runnable object can be executed by a `Thread`
- The `run()` method is executed, when you invoke `start()` on the thread

```
public class MyRunnable implements Runnable {  
    public void run() {  
        System.out.println("MyRunnable running");  
    }  
}
```

- To execute the `run()` method, pass an instance of `Runnable` to a `Thread` in its constructor

```
Thread someThread = new Thread(new MyRunnable());  
someThread.start();
```

- When the thread is started it will call the `run()` method of the `MyRunnable` instance instead of executing it's own `run()` method.

A Common Pitfall

- What is wrong in the code below?

```
public class MyRunnable implements Runnable {
    @Override
    public void run() {
        System.out.println("Runnable running with: " +
            Thread.currentThread().getName());
    }
    public static void main(String[] args) {
        System.out.println(Thread.currentThread().getName());
        Thread aThread = new Thread(new MyRunnable());
        System.out.println(aThread.getName());
        aThread.run(); // Incorrect! Should be aThread.start();
    }
}
```

- It appears that the Runnable's `run()` method executed as expected, but it is NOT executed by the new thread you just created. Instead the `run()` method is executed by the thread that created the thread i.e. the Main thread
- You MUST invoke `someThread.start()` on the thread

A Simple Thread Program

```
public class Counter implements Runnable {

    private static int counter = 0;
    public Counter(int counter) {
        this.counter = counter; }

    @Override
    public void run() {
        for(int i=0; i<5; i++) {
            System.out.println (Thread.currentThread().getName() + ":"
                                + counter++);
        }
    }

    public static void main(String[] args) {
        Runnable r1 = new Counter(counter);
        Thread t1 = new Thread(r1);
        Thread t2 = new Thread(r1);
        t1.start();  t2.start();
    }
}
```

Critical Section & Race Condition

- A **race condition** is a special condition that may occur inside a critical section
- A **critical section** is a section of code that is executed by multiple threads and where the sequence of execution for the threads makes a difference in the result of the concurrent execution of the critical section
- When the result of multiple threads executing a critical section may differ depending on the sequence in which the threads execute, **the critical section is said to contain a race condition**

BoundedQueue class

```
public class BoundedQueue<E>
{
    private Object[] elements;
    private int head;
    private int tail;
    private int size;

    public void add(E newValue) {
        elements[tail] = newValue;
        tail++;
        size++;
        if (tail == elements.length)
            tail = 0;
    }
}
```

```
    public boolean isFull() {
        return size == elements.length;
    }

    public E remove() {
        E r = elements[head];
        head++;
        size--;
        if (head == elements.length)
            head = 0;
        return r;
    }

    public boolean isEmpty() {
        return size == 0;
    }
}
```

Producer Class

```
public class Producer implements Runnable
{
    private String greeting;
    private BoundedQueue<String> queue;
    private int greetingCount;

    public Producer(BoundedQueue<String> aQueue, int count)
    {
        queue = aQueue; greetingCount = count;
    }
    public void run() {
        try {
            int i = 1;
            while (i <= greetingCount) {
                if (!queue.isEmpty()) {
                    String greeting = queue.remove();
                    System.out.println(greeting);
                    i++;
                }
            }
            //... Rest of the code
        }
    }
}
```

Consumer Class

```
public class Consumer implements Runnable
{
    public Consumer(BoundedQueue<String> aQueue, int count)
    {
        queue = aQueue;
        greetingCount = count;
    }

    public void run() {
        try
        {
            int i = 1;
            while (i <= greetingCount) {
                if (!queue.isEmpty()) {
                    String greeting = queue.remove();
                    System.out.println(greeting);
                    i++;
                }
            }
        }
        //... Rest of the code
    }
}
```

ThreadTester class

How could size be 11? (Race-Condition)

```
public class ThreadTester {  
  
    BoundedQueue<String> = new BoundedQueue<String>(10);  
  
    Runnable run1 = new Producer("Hello", queue);  
    Runnable run2 = new Producer("Hello", queue);  
    Runnable run3 = new Consumer("Goodbye", queue);  
  
    Thread t1 = new Thread(run1);  
    Thread t2 = new Thread(run2);  
    Thread t3 = new Thread(run3);  
  
    t1.start();  
    t2.start();  
    t3.start();  
}
```

Lecture Demo: Race Condition

How can the queue be corrupted?

How could size be 11?

Using Java's original synchronization primitives

- Every Java object has an associated **object lock**
- Acquire and release this lock, simply tag the method with the **synchronized** keyword

```
// BoundedQueue class

public synchronized void add(E newValue)
{ ... }

public synchronized E remove()
{ ... }
```

```
// Insider producer

If (!queue. isFull()) {

    queue.add(...);

}
```

Why doesn't this work? – Queue corrupted?

- Solution: the test **isFull()** must be moved to **add()**

```
// BoundedQueue class

public synchronized void add(E newValue)
{ try {
    while (queue is full)
        sleep(); // wait for more space
    ...
}
```

This still doesn't work...

- If the thread sleeps after acquiring lock, no other thread can remove elements
- A consumer thread calls **remove()** but will be blocked until **add()** exits ... **Deadlock!!**

Problems with Locks (1)

A **Deadlocks** arises when locking threads result in a situation where they cannot proceed and thus wait indefinitely for others to terminate e.g.,

- One thread acquires a lock on resource r1 and waits to acquire another on resource r2 and at the same time, another thread that has already acquired r2 is waiting to obtain a lock on r1
- Here neither thread can proceed until the other one releases the lock, which never happens – resulting in a deadlock

Problems with Locks (2)

- A **Lock Starvation** arises when threads have different priorities and a thread scheduler gives lock to high-priority threads.
- If there are many high-priority threads that want to obtain the lock and also hold the lock for long time periods, this could result in a situation where low-priority threads “starve” for a long time trying to obtain the lock

Revised bounded queue

```
public synchronized void
    add(E newValue) {
    while (isFull()) wait();
    elements[tail] = newValue;
    tail++;
    size++;
    if (tail ==
elements.length)
        tail = 0;
    notifyAll();
}
```

```
public synchronized E
    remove() {
    while (isEmpty()) wait();
    E r = elements[head];
    head++;
    size--;
    if (head == elements.length)
        head = 0;
    notifyAll();
}
```

- This does work, but can it be better?

Synchronisation using Object Locks

```
public void add(E newValue) {  
    queueLock.lock() ;  
    while (isFull()) sleep() ;  
    elements[tail] = newValue;  
    tail++;  
    size++;  
    if (tail == elements.length)  
        tail = 0;  
    queueLock.unlock() ;  
}
```

```
public E remove() {  
    queueLock.lock() ;  
    while (isEmpty()) sleep() ;  
    E r = elements[head];  
    head++;  
    size--;  
    if (head ==  
        elements.length)  
        head = 0;  
    queueLock.unlock() ;  
}
```

Java's new **Lock** interface, is preferred over using primitive **synchronised**...but the code above has still same issue
deadlock

Reentrant Locks: Using Lock and Condition Objects

```
private Lock queueLock = new ReentrantLock();  
private Condition spaceAvailableCond = queueLock.newCondition();  
private Condition valueAvailableCond = queueLock.newCondition();
```

```
public void add(E newValue) {  
    queueLock.lock();  
    try {  
        while (isFull())  
            space.await();  
        elements[tail] = newValue;  
        tail++;  
        size++;  
        if (tail == elements.length)  
            tail = 0;  
        valueAvailableCond.signalAll();  
    }  
    finally {queueLock.unlock();}  
}
```

```
public E remove() {  
    queueLock.lock();  
    try {  
        while (isEmpty())  
            value.await();  
        E r = elements[head];  
        head++;  
        size--;  
        if (head == elements.length)  
            head = 0;  
        spaceAvailableCond.signalAll();  
    }  
    finally { queueLock.unlock();}  
}
```